

```

import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
import seaborn as sns

import statsmodels.api as sm

import math

import datetime

import warnings
warnings.filterwarnings('ignore')

from math import sqrt

from scipy import stats

df = pd.read_csv("/content/WalmartDataSet.csv")

df.shape

(6435, 8)

df.columns

Index(['Store', 'Date', 'Weekly_Sales', 'Holiday_Flag', 'Temperature',
       'Fuel_Price', 'CPI', 'Unemployment'],
      dtype='object')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Store           6435 non-null   int64
 1   Date            6435 non-null   object
 2   Weekly_Sales    6435 non-null   float64
 3   Holiday_Flag    6435 non-null   int64
 4   Temperature     6435 non-null   float64
 5   Fuel_Price      6435 non-null   float64
 6   CPI             6435 non-null   float64
 7   Unemployment    6435 non-null   float64
dtypes: float64(5), int64(2), object(1)
memory usage: 402.3+ KB

# Change Store column to categorical type
df['Store'] = df.Store.astype('category')

```

```
# Change Holiday_Flag column to boolean type
df['Holiday_Flag'] = df['Holiday_Flag'].astype(bool)
```

```
# Change Date columns to datetime type
df["Date"] = pd.to_datetime(df["Date"])
```

```
df['Year'] =df['Date'].dt.year
df['Month'] =df['Date'].dt.month
df['Week'] =df['Date'].dt.week
```

```
df.head()
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature
0	1	2010-05-02	1643690.90	False	42.31
1	1	2010-12-02	1641957.44	True	38.51
2	1	2010-02-19	1611968.17	False	39.93
3	1	2010-02-26	1409727.59	False	46.63
4	1	2010-05-03	1554806.68	False	46.50

	CPI	Unemployment	Year	Month	Week
0	211.096358	8.106	2010	5	17
1	211.242170	8.106	2010	12	48
2	211.289143	8.106	2010	2	7
3	211.319643	8.106	2010	2	8
4	211.350143	8.106	2010	5	18

```
df.tail()
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature
6430	45	2012-09-28	713173.95	False	64.88
6431	45	2012-05-10	733455.07	False	64.89
6432	45	2012-12-10	734464.36	False	54.47
6433	45	2012-10-19	718125.53	False	56.47
6434	45	2012-10-26	760281.43	False	58.85

	CPI	Unemployment	Year	Month	Week
6430	192.013558	8.684	2012	9	39
6431	192.170412	8.667	2012	5	19
6432	192.327265	8.667	2012	12	50

6433	192.330854	8.667	2012	10	42
6434	192.308899	8.667	2012	10	43

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 6435 entries, 0 to 6434
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	Store	6435 non-null	category
1	Date	6435 non-null	datetime64[ns]
2	Weekly_Sales	6435 non-null	float64
3	Holiday_Flag	6435 non-null	bool
4	Temperature	6435 non-null	float64
5	Fuel_Price	6435 non-null	float64
6	CPI	6435 non-null	float64
7	Unemployment	6435 non-null	float64
8	Year	6435 non-null	int64
9	Month	6435 non-null	int64
10	Week	6435 non-null	int64

```
dtypes: bool(1), category(1), datetime64[ns](1), float64(5), int64(3)
```

```
memory usage: 466.6 KB
```

```
df.isnull().sum()
```

Store	0
Date	0
Weekly_Sales	0
Holiday_Flag	0
Temperature	0
Fuel_Price	0
CPI	0
Unemployment	0
Year	0
Month	0
Week	0

```
dtype: int64
```

```
df.isnull()
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature
Fuel_Price \					
0	False	False	False	False	False
False					
1	False	False	False	False	False
False					
2	False	False	False	False	False
False					
3	False	False	False	False	False
False					

```

4      False  False      False      False      False
False
...      ...      ...      ...      ...      ...
.
6430   False  False      False      False      False
False
6431   False  False      False      False      False
False
6432   False  False      False      False      False
False
6433   False  False      False      False      False
False
6434   False  False      False      False      False
False

```

```

      CPI  Unemployment  Year  Month  Week
0      False      False  False  False  False
1      False      False  False  False  False
2      False      False  False  False  False
3      False      False  False  False  False
4      False      False  False  False  False
...      ...      ...      ...      ...
6430   False      False  False  False  False
6431   False      False  False  False  False
6432   False      False  False  False  False
6433   False      False  False  False  False
6434   False      False  False  False  False

```

```
[6435 rows x 11 columns]
```

```

# Exclude datetime column(s) from mean imputation
mean_imputation_cols = df.columns.drop(['Date'])

# Fill missing values with mean imputation
df.fillna(df.mean(), inplace=True)

# Check the number of missing values in each column
df.isnull().sum()

```

```

Store      0
Date        0
Weekly_Sales  0
Holiday_Flag  0
Temperature  0
Fuel_Price  0
CPI         0
Unemployment 0
Year        0
Month       0

```

```

Week          0
dtype: int64

# Check for duplicate rows
df.duplicated().sum() # Count the number of duplicate rows

0

# Remove duplicate rows
df.drop_duplicates(inplace=True)

# Check for duplicate columns
df.columns.duplicated().sum() # Count the number of duplicate columns

0

# Remove duplicate columns, if any
df = df.loc[:, ~df.columns.duplicated()]

Q1 = df['Weekly_Sales'].quantile(0.25)
Q3 = df['Weekly_Sales'].quantile(0.75)
IQR = Q3 - Q1
outlier_threshold = 1.5 * IQR

df = df[(df['Weekly_Sales'] >= Q1 - outlier_threshold) &
(df['Weekly_Sales'] <= Q3 + outlier_threshold)]

```

## EDA

```
df.describe()
```

	Weekly_Sales	Temperature	Fuel_Price	CPI
Unemployment \				
count	6.401000e+03	6401.000000	6401.000000	6401.000000
mean	1.036130e+06	60.772042	3.359634	171.642219
std	5.451961e+05	18.417068	0.459696	39.359852
min	2.099862e+05	-2.060000	2.472000	126.064000
25%	5.517431e+05	47.660000	2.933000	131.784000
50%	9.572983e+05	62.860000	3.452000	182.658578
75%	1.414565e+06	75.000000	3.737000	212.833640
max	2.685352e+06	100.140000	4.468000	227.232807

	Year	Month	Week
count	6401.000000	6401.000000	6401.000000
mean	2010.967974	6.447899	25.875644
std	0.797304	3.308627	14.448448
min	2010.000000	1.000000	1.000000
25%	2010.000000	4.000000	14.000000
50%	2011.000000	6.000000	26.000000
75%	2012.000000	9.000000	38.000000
max	2012.000000	12.000000	52.000000

Min Date & Max Date

```
print('Min Date in Data is - {}'.format(df['Date'].min()))
print('Max Date in Data is - {}'.format(df['Date'].max()))
```

```
Min Date in Data is - 2010-01-10 00:00:00
Max Date in Data is - 2012-12-10 00:00:00
```

```
# Scatter plot to explore the relationship between weekly sales and unemployment rate
```

```
plt.figure(figsize=(20, 5))
sns.scatterplot(x='Weekly_Sales', y='Unemployment', data=df)
```

```
# Set the title
```

```
plt.title('Correlation between Weekly Sales and Unemployment',
color='#007DC6', fontsize=20, pad=10)
```

```
# Set the x and y axis labels
```


```
plt.xlabel('Weekly Sales', color='#F1C21B', fontweight='bold',
fontsize=14)
plt.ylabel('Unemployment', color='#F1C21B', fontweight='bold',
fontsize=14)
```

```
# Adjust plot layout
```

```
plt.tight_layout()
```

```
# Display the plot
```

```
plt.show()
```



```
# Calculate the correlation coefficient between weekly sales and unemployment rate
correlation_coefficient = df['Weekly_Sales'].corr(df['Unemployment'])
print("Correlation Coefficient Between Weekly Sales and Unemployment Rate:", correlation_coefficient)
```

Correlation Coefficient Between Weekly Sales and Unemployment Rate: -0.10429750912578388

```
# Calculate the correlation between each store's weekly sales and unemployment rate
```

```
store_correlation = df.groupby('Store')['Weekly_Sales', 'Unemployment'].corr().unstack()
```

```
# Find stores with the highest negative correlation
```

```
stores_with_highest_negative_correlation = store_correlation['Weekly_Sales']['Unemployment'].idxmin()
```

```
print("Stores with the Highest Negative Correlation with Unemployment Rate:", stores_with_highest_negative_correlation)
```

Stores with the Highest Negative Correlation with Unemployment Rate: 38

```
# Find the top N number of stores with the highest negative correlation
```

```
n = 10 # specify the number of stores to display, you can modify this value
```

```
print("Top", n, "Stores with the Highest Negative Correlation with Unemployment Rate:")
```

```
print()
```

```
top_negative_corr_stores = store_correlation['Weekly_Sales']['Unemployment'].sort_values().head(n)
```

```
# Create a dataframe
```

```
top_negative_corr_stores_df = pd.DataFrame({'Store': top_negative_corr_stores.index, 'Negative_Correlation': top_negative_corr_stores.values})
```

```
# Set the index to start from 1
```

```
top_negative_corr_stores_df.index = top_negative_corr_stores_df.index + 1
```

```
print(top_negative_corr_stores_df.to_string(index=False))
```

Top 10 Stores with the Highest Negative Correlation with Unemployment Rate:

Store	Negative_Correlation
38	-0.785290
44	-0.780076
4	-0.639563
13	-0.400254
39	-0.384681
42	-0.356355
41	-0.350630
17	-0.263600
3	-0.230413
37	-0.221287

---


How Holidays affect the sales of each store.

```
# impact of holidays on weekly sales
plt.figure(figsize = (12,8))
sns.boxplot(x = 'Holiday_Flag', y = 'Weekly_Sales', data = df,
showliers = False)

# Mapping 0 to False and 1 to True in x-axis labels
plt.xticks([0, 1], ['False', 'True'])


plt.show()
```





## Data Visualization

```
# Analyzing the distribution of target variable
plt.figure(figsize = (10, 5))
sns.distplot(df['Weekly_Sales'], hist_kws=dict(edgecolor="black"))
plt.title('Weekly Sales Distribution', fontsize= 15)
plt.grid()
plt.show()
```



Histogram of Weekly Sales

```
# Set the size of the graph
plt.figure(figsize=(10, 6))

# Plot a histogram of the 'Weekly_Sales' column from the DataFrame
plt.hist(df['Weekly_Sales'], bins=20)

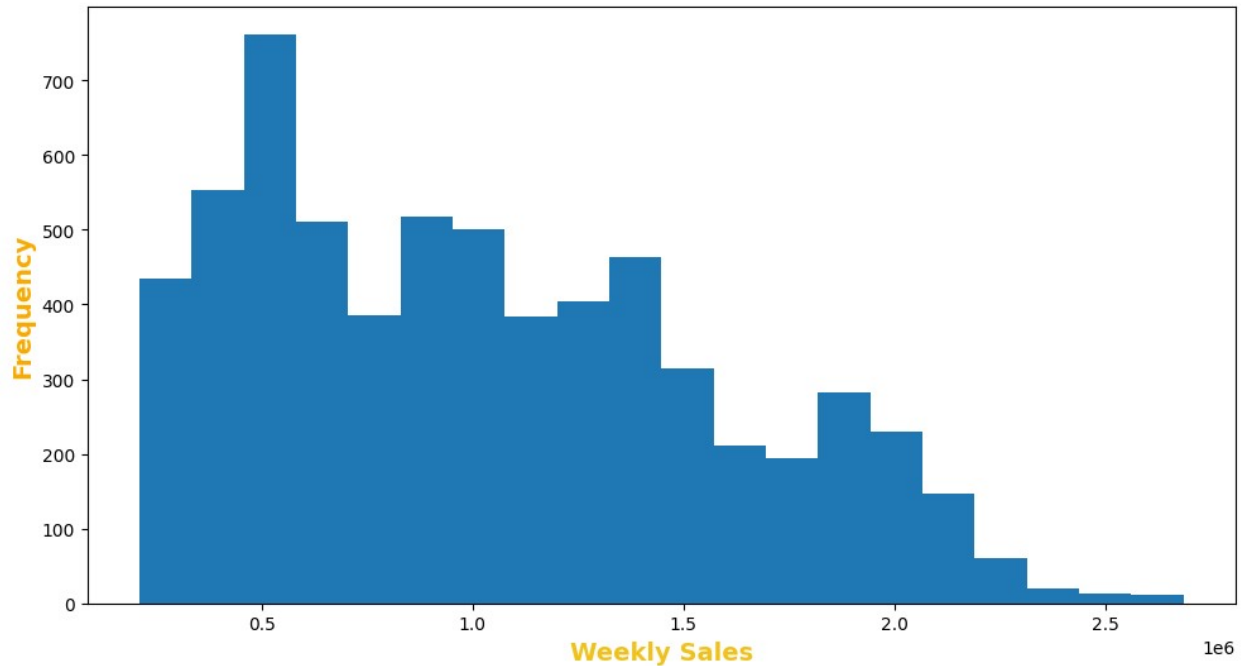
# Set the x and y labels with specified colors, font weight, and font size
plt.xlabel('Weekly Sales', color='#F1C21B', fontweight='bold',
           fontsize=14)
plt.ylabel('Frequency', color='#F9AB00', fontweight='bold',
           fontsize=14)

# Set the title with specified color, font style, font size, and padding
plt.title('Distribution of Weekly Sales', color='#007DC6',
          fontsize=20, pad=20)

# Adjust plot appearance and layout
plt.tight_layout()

# Display the plot
plt.show()
```

Distribution of Weekly Sales



```
df.groupby('Month')['Weekly_Sales'].mean()
```

Month

1	9.476139e+05
2	1.054597e+06
3	1.024975e+06
4	1.024324e+06
5	1.035379e+06
6	1.064848e+06
7	1.014212e+06
8	1.044874e+06
9	1.009457e+06
10	1.027683e+06
11	1.093977e+06
12	1.110051e+06

Name: Weekly\_Sales, dtype: float64

```
df.groupby('Year')['Weekly_Sales'].mean()
```

Year

2010	1.040919e+06
2011	1.033780e+06
2012	1.033660e+06

Name: Weekly\_Sales, dtype: float64

```
df['Holiday_Flag'].value_counts()
```

```
False    5960
True      441
Name: Holiday_Flag, dtype: int64
```

Relation between a categorized feature and the Weekly\_Sales

```
def graph_relation_to_weekly_sale(col_relation, df, x='Week',
palette=None):
    cmap = sns.diverging_palette(220, 20, as_cmap=True)

    sns.relplot(
        x=x,
        y='Weekly_Sales',
        hue=col_relation,
        data=df,
        kind='line',
        height=5,
        aspect=2,
        palette=palette
    )
    plt.show()


graph_relation_to_weekly_sale('Year', df, x='Date', palette='Set2')
```



Average Weekly Sales

```
plt.figure(figsize=(20, 7))
sns.barplot(x='Week', y='Weekly_Sales', data=df)
plt.title('Average Weekly Sales', fontsize=18)
plt.ylabel('Weekly Sales', fontsize=16)
plt.xlabel('Week', fontsize=16)
```

```
plt.grid()
plt.show()
```



Box plot of Weekly Sales by Store

```
plt.figure(figsize=(20, 8))

# Plot the boxplot
sns.boxplot(x='Store', y='Weekly_Sales', data=df)


# Set the x and y labels with Walmart font colors
plt.xlabel('Store', color='#007DC6', fontweight='bold', fontsize=20)
plt.ylabel('Weekly Sales', color='#F1C21B', fontweight='bold',
           fontsize=20)

# Set the title with Walmart font colors
plt.title('Distribution of Weekly Sales by Store', color='#2FB12C',
          fontsize=30, pad=20)

# Set the font size for x-axis labels
plt.xticks(fontsize=10)


# Adjust plot layout
plt.tight_layout()

# Display the plot
plt.show()
```



## Heatmap

```
plt.figure(figsize=(20, 10))
sns.heatmap(df.corr(), cmap='PuBu', annot=True)
plt.title('Correlation Matrix', fontsize=16)
plt.show()
```



## Correlation Graph

```
# Correlation between Unemployment & Store
plt.figure(figsize=(20, 5))
sns.set_theme(style="whitegrid")
```

```

ax = sns.barplot(x='Store', y="Unemployment", data=df)


# Set the title
plt.title('Correlation Between Unemployment & Store', color='#007DC6',
          fontsize=20, pad=10)

# Set the x and y axis labels
plt.xlabel('Store', color='#F1C21B', fontweight='bold', fontsize=14)
plt.ylabel('Weekly Sales', color='#F1C21B', fontweight='bold',
           fontsize=14)

# Adjust plot layout
plt.tight_layout()

# Display the plot
plt.show()

```



```

# Correlation between Weekly Sales & Store
plt.figure(figsize=(20, 5))
sns.set_theme(style="whitegrid")
ax = sns.barplot(x='Store', y="Weekly_Sales", data=df)


# Set the title
plt.title('Correlation Between Weekly Sales & Store', color='#007DC6',
          fontsize=20, pad=10)

# Set the x and y axis labels
plt.xlabel('Store', color='#F1C21B', fontweight='bold', fontsize=14)
plt.ylabel('Weekly Sales', color='#F1C21B', fontweight='bold',
           fontsize=14)

# Adjust plot layout
plt.tight_layout()

# Display the plot
plt.show()

```




```
# Correlation b/w Fuel_Price & Year
plt.figure(figsize=(20, 5))
sns.set_theme(style="whitegrid")
ax = sns.barplot(x="Year", y="Fuel_Price", data=df)
sns.set(rc = {'figure.figsize': (10,4)})

# Set the title
plt.title('Correlation Between Fuel Price & Year', color='#007DC6',
fontsize=20, pad=10)

# Set the x and y axis labels
plt.xlabel('Year', color='#F1C21B', fontweight='bold', fontsize=14)
plt.ylabel('Fuel Price', color='#F1C21B',
fontweight='bold', fontsize=14)

# Adjust plot layout
plt.tight_layout()

# Display the plot
plt.show()
```



```
# Correlation between Weekly Sales & Month
plt.figure(figsize=(20, 5))
sns.set_theme(style="whitegrid")
month_wise_sales = df.pivot_table(values="Weekly_Sales",
columns="Year", index="Month")
month_wise_sales.plot(marker='o')

# Set the title with Walmart font colors
plt.title('Correlation Between Weekly Sales and Month',
```



```

color='#007DC6', fontsize=16, pad=10)


# Set the x and y labels with Walmart font colors
plt.xlabel('Month', color='#F1C21B', fontweight='bold', fontsize=14)
plt.ylabel('Weekly Sales', color='#F1C21B', fontweight='bold',
fontsize=14)

# Adjust plot layout
plt.legend(loc='upper left')
plt.tight_layout()

# Display the plot
plt.show()

<Figure size 2000x500 with 0 Axes>

```



```

# Correlation between Unemployment & Store
plt.figure(figsize=(20, 5))
sns.set_theme(style="whitegrid")

# Draw Scatterplot
sns.scatterplot(x="Store", y="Unemployment", data=df)


# Set the title
plt.title('Correlation Between Unemployment & Store', color='#007DC6',
fontsize=20, pad=10)

# Set the x and y axis labels
plt.xlabel('Store', color='#F1C21B', fontweight='bold', fontsize=14)
plt.ylabel('Unemployment', color='#F1C21B', fontweight='bold',
fontsize=14)

# Adjust plot layout
plt.tight_layout()

```

```
# Display the plot
plt.show()
```



```
# Scatter plot of Weekly Sales vs. Temperature
plt.figure(figsize=(10, 6))
```

```
# Plot the scatter plot
sns.scatterplot(x="Temperature", y="Weekly_Sales", data=df)
```


```
# Set the title with Walmart font colors
plt.title('Correlation Between Weekly Sales & Temperature',
color='#007DC6', fontsize=16, pad=10)
```

```
# Set the x and y labels with Walmart font colors
plt.xlabel('Temperature', color='#F1C21B', fontweight='bold',
fontsize=14)
plt.ylabel('Weekly Sales', color='#F1C21B', fontweight='bold',
fontsize=14)
```

```
# Set the font size for x-axis labels
plt.xticks(fontsize=10)
```

```
# Adjust plot layout
plt.tight_layout()
```

```
# Display the plot
plt.show()
```



```
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf

from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor

from xgboost import XGBRegressor

from sklearn import tree
from sklearn.preprocessing import OneHotEncoder

from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.model_selection import train_test_split

from sklearn.metrics import *
from sklearn import metrics
```

---

```
df.columns
```

```
Index(['Store', 'Date', 'Weekly_Sales', 'Holiday_Flag', 'Temperature',
       'Fuel_Price', 'CPI', 'Unemployment', 'Year', 'Month', 'Week'],
      dtype='object')
```

```
x = df.drop(['Date', 'Weekly_Sales'], axis=1)
```

```
x
```

	Store Unemployment	Holiday_Flag \	Temperature	Fuel_Price	CPI
0	1	False	42.31	2.572	211.096358
8.106					
1	1	True	38.51	2.548	211.242170
8.106					
2	1	False	39.93	2.514	211.289143
8.106					
3	1	False	46.63	2.561	211.319643
8.106					
4	1	False	46.50	2.625	211.350143
8.106					
...	...	...	...	...	...
...					
6430	45	False	64.88	3.997	192.013558
8.684					
6431	45	False	64.89	3.985	192.170412
8.667					
6432	45	False	54.47	4.000	192.327265
8.667					
6433	45	False	56.47	3.969	192.330854
8.667					
6434	45	False	58.85	3.882	192.308899
8.667					

	Year	Month	Week
0	2010	5	17
1	2010	12	48
2	2010	2	7
3	2010	2	8
4	2010	5	18
...	...	...	...
6430	2012	9	39
6431	2012	5	19
6432	2012	12	50
6433	2012	10	42
6434	2012	10	43

```
[6401 rows x 9 columns]
```

```
x.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6401 entries, 0 to 6434
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype

```

```

---
0  Store          6401 non-null  category
1  Holiday_Flag  6401 non-null  bool
2  Temperature   6401 non-null  float64
3  Fuel_Price    6401 non-null  float64
4  CPI           6401 non-null  float64
5  Unemployment  6401 non-null  float64
6  Year          6401 non-null  int64
7  Month         6401 non-null  int64
8  Week          6401 non-null  int64
dtypes: bool(1), category(1), float64(4), int64(3)
memory usage: 672.0 KB

```


```
y = df['Weekly_Sales']
```

```
rf = RandomForestRegressor(n_estimators = 100)
rf.fit(x, y)
```

```
RandomForestRegressor()
```

```
# checking the feature importance
```

```
plt.figure(figsize = (15, 5))
plt.bar(x.columns, rf.feature_importances_)
plt.title("Feature Importance", fontsize = 15)
plt.show()
```



```
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size = 0.8, random_state = 0)
```

```
#Linear Regression
```

```
lr = LinearRegression()
```

```
lr.fit(x_train, y_train)
```


```
LinearRegression()
```

```
y_pred = lr.predict(x_test)
```

```
plt.scatter(y_test, y_pred)
```

```
print("R2 Score: ", r2_score(y_test, y_pred))
print("MSE Score: ", mean_squared_error(y_test, y_pred))
print("RMSE : ", sqrt(mean_squared_error(y_test, y_pred)))
```


R2 Score: 0.1640905885815509  
MSE Score: 270519048713.93335  
RMSE : 520114.4573206299



```
#Decision Tree
dtree = DecisionTreeRegressor()
dtree.fit(x_train, y_train)
DecisionTreeRegressor()
y_pred1 = dtree.predict(x_test)
plt.scatter(y_test, y_pred1)

print("R2 Score: ", r2_score(y_test, y_pred1))
print("MSE Score: ", mean_squared_error(y_test, y_pred1))
print("RMSE : ", sqrt(mean_squared_error(y_test, y_pred1)))
```


R2 Score: 0.9203742962187489  
MSE Score: 25768665056.097343  
RMSE : 160526.21298746613



```
#Random Forest
rf1 = RandomForestRegressor(n_estimators = 100)
rf1.fit(x_train, y_train)
RandomForestRegressor()
y_pred2 = rf1.predict(x_test)
plt.scatter(y_test, y_pred2)

print("R2 Score: ", r2_score(y_test, y_pred2))
print("MSE Score: ", mean_squared_error(y_test, y_pred2))
print("RMSE : ", sqrt(mean_squared_error(y_test, y_pred2)))
```


```
R2 Score: 0.9530368529163145
MSE Score: 15198328551.096323
RMSE : 123281.50125260612
```



*#KNN*

```
knn = KNeighborsRegressor()  
knn.fit(x_train, y_train)  
KNeighborsRegressor()  
y_pred3 = knn.predict(x_test)  
plt.scatter(y_test, y_pred3)  
  
print("R2 Score: ", r2_score(y_test, y_pred3))  
print("MSE Score: ", mean_squared_error(y_test, y_pred3))  
print("RMSE : ", sqrt(mean_squared_error(y_test, y_pred3)))
```

```
R2 Score:  0.4890754165027007  
MSE Score: 165346663650.68936  
RMSE :  406628.4097928837
```



*#XGBoost*

*# Specify the features and target variable*

```
features = ['Store', 'Year', 'Month']  
target = 'Weekly_Sales'
```

*# Split the data into training and testing sets*

```
x_train, x_test, y_train, y_test = train_test_split(df[features],  
df[target], test_size=0.2, random_state=42)
```

*# Create and train the XGBoost model with enable\_categorical=True*

```
xg = XGBRegressor(n_estimators=100, random_state=42,  
enable_categorical=True)  
xg.fit(x_train, y_train)
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,  
colsample_bylevel=None, colsample_bynode=None,
```



```
        colsample_bytree=None, device=None,
early_stopping_rounds=None,
        enable_categorical=True, eval_metric=None,
feature_types=None,
        gamma=None, grow_policy=None, importance_type=None,
        interaction_constraints=None, learning_rate=None,
max_bin=None,
        max_cat_threshold=None, max_cat_to_onehot=None,
        max_delta_step=None, max_depth=None, max_leaves=None,
        min_child_weight=None, missing=nan,
monotone_constraints=None,
        multi_strategy=None, n_estimators=100, n_jobs=None,
        num_parallel_tree=None, random_state=42, ...)
```

```
# Make predictions on the test set
```

```
y_pred4 = xg.predict(x_test)
```

```
# Evaluate the model
```

```
r2 = r2_score(y_test, y_pred4)
```

```
mse = mean_squared_error(y_test, y_pred4)
```

```
rmse = np.sqrt(mse)
```

```
print("R2 Score: ", r2)
```

```
print("MSE Score: ", mse)
```

```
print("RMSE Score: ", rmse)
```

```
# Plot predicted vs actual values
```

```
plt.scatter(y_test, y_pred4)
```

```
plt.xlabel("Actual Weekly Sales")
```

```
plt.ylabel("Predicted Weekly Sales")
```


```
plt.title("XGBoost: Predicted vs Actual Weekly Sales")
```

```
plt.show()
```

```
R2 Score: 0.9248258065451529
```

```
MSE Score: 23872625335.60029
```


```
RMSE Score: 154507.68697899886
```



```
# Getting Average of Best Models
y_pred_final = (y_pred1 + y_pred2 + y_pred4)/3.0
plt.scatter(y_test, y_pred_final)

print("R2 Score: ", r2_score(y_test, y_pred_final))
print("MSE Score: ", mean_squared_error(y_test, y_pred_final))
print("RMSE : ", sqrt(mean_squared_error(y_test, y_pred_final)))

R2 Score: 0.08903920858824954
MSE Score: 289288446863.84094
RMSE : 537855.4144599094
```



```

# Calculate summary statistics
summary_stats = df.describe()

# Perform statistical tests or correlations
# Example: Calculate correlation between Weekly Sales and Unemployment Rate
correlation = df['Weekly_Sales'].corr(df['Unemployment'])

# Additional statistical tests
# Example: t-test comparing Weekly Sales between two groups (Group_A and Group_B)
group_a_sales = df[df['Holiday_Flag'] == 0]['Weekly_Sales']
group_b_sales = df[df['Holiday_Flag'] == 1]['Weekly_Sales']
t_stat, p_value = stats.ttest_ind(group_a_sales, group_b_sales)

```

Perform t-test

```

# Perform t-test
t_stat, p_value = stats.ttest_ind(group_a_sales, group_b_sales)

# Print the results
print("T-Statistic:", t_stat)
print("P-Value:", p_value)

T-Statistic: -2.029118506966506
P-Value: 0.042487534639960314

```

#OR

```

# Perform t-test
#t_stat, p_value = stats.ttest_ind(group_a_sales, group_b_sales)

# Create a DataFrame to display the results
#results_df = pd.DataFrame({'T-Statistic': [t_stat], 'P-Value': [p_value]})

# Display the results
#display(results_df)

#3. Apply time series analysis to identify seasonal trends:




# Use seasonal decomposition to identify seasonal, trend, and residual components
result = sm.tsa.seasonal_decompose(df['Weekly_Sales'],
model='additive', period=52)

# Plot the decomposed components
fig, ax = plt.subplots(4, 1, figsize=(20, 40))
result.observed.plot(ax=ax[0])
ax[0].set_ylabel('Observed')

```

```
result.trend.plot(ax=ax[1])
ax[1].set_ylabel('Trend')
result.seasonal.plot(ax=ax[2])
ax[2].set_ylabel('Seasonal')
result.resid.plot(ax=ax[3])
ax[3].set_ylabel('Residual')

plt.tight_layout()
plt.show()
```



```
#4. Analyze the impact of variables on weekly sales:

# Assuming you have additional variables (unemployment rate,
temperature, and CPI) in your dataset

# Create a matrix of independent variables
X = df[['Unemployment', 'Temperature', 'CPI']]

# Add constant to the matrix
X = sm.add_constant(X)

# Fit the multiple linear regression model
model = sm.OLS(df['Weekly_Sales'], X)
results = model.fit()

# Print the model summary
print(results.summary())
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:          Weekly_Sales    R-squared:
0.022
Model:                  OLS            Adj. R-squared:
0.022
Method:                 Least Squares   F-statistic:
48.68
Date:                   Tue, 21 Nov 2023 Prob (F-statistic):
4.22e-31
Time:                   10:31:55        Log-Likelihood:
-93560.
No. Observations:      6401            AIC:
1.871e+05
Df Residuals:          6397            BIC:
1.872e+05
Df Model:               3

Covariance Type:       nonrobust

=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	1.633e+06	5.09e+04	32.065	0.000	1.53e+06	1.73e+06
Unemployment	-3.956e+04	3817.156	-10.363	0.000	-4.7e+04	-3.21e+04

Temperature	-342.2836	376.853	-0.908	0.364	-1081.041
396.474					
CPI	-1510.3415	184.191	-8.200	0.000	-1871.417
-1149.266					

```
=====
=====
Omnibus:                    554.908    Durbin-Watson:
0.090
Prob(Omnibus):              0.000    Jarque-Bera (JB):
371.066
Skew:                      0.474    Prob(JB):
2.65e-81
Kurtosis:                  2.299    Cond. No.
1.41e+03
=====
=====
```


#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.41e+03. This might indicate that there are strong multicollinearity or other numerical problems.

#### #5. Visualize relationships using appropriate plots:

```
# Scatter plots of the variables against weekly sales
fig, axs = plt.subplots(3, 1, figsize=(20, 10))
axs[0].scatter(df['Unemployment'], df['Weekly_Sales'])
axs[0].set_xlabel('Unemployment Rate')
axs[0].set_ylabel('Weekly Sales')
axs[1].scatter(df['Temperature'], df['Weekly_Sales'])
axs[1].set_xlabel('Temperature')
axs[1].set_ylabel('Weekly Sales')
axs[2].scatter(df['CPI'], df['Weekly_Sales'])
axs[2].set_xlabel('CPI')
axs[2].set_ylabel('Weekly Sales')

plt.tight_layout()
plt.show()
```



*#2. Aggregate the weekly sales data by store and calculate total sales:*

*# Assuming you have already loaded and preprocessed the Walmart dataset*

*# Group the dataset by store*

```
df_by_store = df.groupby('Store')['Weekly_Sales'].sum().reset_index()
```

*# Sort the stores based on total sales*

```
df_by_store_sorted = df_by_store.sort_values('Weekly_Sales',
ascending=False)
```

*# Print the top 5 best performing stores*

```
print("Top 5 Best Performing Stores:")
```

```
print(df_by_store_sorted.head())
```

*# Print Empty Row*

```
print("")
```

*# Print the worst performing stores*

```
print("Top 5 Worst Performing Store:")
```

```
print(df_by_store_sorted.tail(5))
```

Top 5 Best Performing Stores:

	Store	Weekly_Sales
3	4	2.810352e+08
19	20	2.800237e+08
13	14	2.761276e+08
1	2	2.687221e+08
12	13	2.682025e+08



Top 5 Worst Performing Store:

	Store	Weekly_Sales
37	38	55159626.42
35	36	53412214.97
4	5	45475688.90
43	44	43293087.84
32	33	37160221.96

*#3. Rank the stores based on their historical sales data to identify the top performing stores:*

*# Add a rank column based on total sales*

```
df_by_store_sorted['Rank'] = np.arange(1, len(df_by_store_sorted) + 1)
```

*# Print the ranked stores*

```
print("Ranked Stores:")
```

```
print(df_by_store_sorted)
```

Ranked Stores:

	Store	Weekly_Sales	Rank
3	4	2.810352e+08	1
19	20	2.800237e+08	2
13	14	2.761276e+08	3
1	2	2.687221e+08	4
12	13	2.682025e+08	5
9	10	2.556789e+08	6
26	27	2.480387e+08	7
0	1	2.224028e+08	8
5	6	2.210286e+08	9
38	39	2.074455e+08	10
18	19	2.066349e+08	11
30	31	1.996139e+08	12
22	23	1.960163e+08	13
23	24	1.940160e+08	14
10	11	1.939628e+08	15
27	28	1.892637e+08	16
40	41	1.813419e+08	17
31	32	1.668192e+08	18
17	18	1.551147e+08	19
21	22	1.470756e+08	20
11	12	1.442872e+08	21
25	26	1.434164e+08	22
33	34	1.382498e+08	23
39	40	1.378703e+08	24
34	35	1.315207e+08	25
7	8	1.299512e+08	26
16	17	1.277821e+08	27
44	45	1.123953e+08	28
20	21	1.081179e+08	29

24	25	1.010612e+08	30
42	43	9.056544e+07	31
14	15	8.913368e+07	32
6	7	8.159828e+07	33
41	42	7.956575e+07	34
8	9	7.778922e+07	35
28	29	7.714155e+07	36
15	16	7.425243e+07	37
36	37	7.420274e+07	38
29	30	6.271689e+07	39
2	3	5.758674e+07	40
37	38	5.515963e+07	41
35	36	5.341221e+07	42
4	5	4.547569e+07	43
43	44	4.329309e+07	44
32	33	3.716022e+07	45

```
# Create a new column 'Date' based on the index
```

```
df['Date'] = df.index
```

```
# If the index is not in the correct date format, convert it to
datetime format
```

```
df['Date'] = pd.to_datetime(df['Date'])
```

```
# Reset the index since we now have the 'Date' column
```

```
df.reset_index(drop=True, inplace=True)
```

```
# Now, the 'Date' column should be restored in your DataFrame
```

```
print(df.head())
```

	Store	Date	Weekly_Sales	Holiday_Flag	\
0	1	1970-01-01 00:00:00.000000000	1643690.90	False	
1	1	1970-01-01 00:00:00.000000001	1641957.44	True	
2	1	1970-01-01 00:00:00.000000002	1611968.17	False	
3	1	1970-01-01 00:00:00.000000003	1409727.59	False	
4	1	1970-01-01 00:00:00.000000004	1554806.68	False	

	Temperature	Fuel_Price	CPI	Unemployment	Year	Month
Week						
0	42.31	2.572	211.096358	8.106	2010	5
17						
1	38.51	2.548	211.242170	8.106	2010	12
48						
2	39.93	2.514	211.289143	8.106	2010	2
7						
3	46.63	2.561	211.319643	8.106	2010	2
8						
4	46.50	2.625	211.350143	8.106	2010	5
18						

```

df.columns

Index(['Store', 'Date', 'Weekly_Sales', 'Holiday_Flag', 'Temperature',
      'Fuel_Price', 'CPI', 'Unemployment', 'Year', 'Month', 'Week'],
      dtype='object')

#2. Load the dataset and preprocess the data:

# Convert the 'Date' column to datetime
df['Date'] = pd.to_datetime(df['Date'])

# Set the 'Date' column as the index
df.set_index('Date', inplace=True)

#3. Split the dataset into training and testing sets:

# Split the dataset into 80% train and 20% test
train_size = int(len(df) * 0.8)
train_data, test_data = df[:train_size], df[train_size:]

#4. Choose a predictive modeling technique:
#- For linear regression:

model = RandomForestRegressor()

#- For time series forecasting, you may use SARIMAX from the
statsmodels library. However, this requires additional preprocessing
and stationarity checks for the time series.

#5. Implement and train the model:

# Assuming you have selected the predictive modeling technique (e.g.,
Random Forest)

# Separate the target variable from the features
X_train, y_train = train_data.drop('Weekly_Sales', axis=1),
train_data['Weekly_Sales']
X_test, y_test = test_data.drop('Weekly_Sales', axis=1),
test_data['Weekly_Sales']

# Fit the model on the training data
model.fit(X_train, y_train)

RandomForestRegressor()

#6. Evaluate the model's performance:

plt.figure(figsize=(20, 8))

# Make predictions on the test set
y_pred = model.predict(X_test)

```

```

# Evaluate the model using mean squared error (MSE)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)

# Visualize the actual vs. predicted values
plt.plot(y_test.index, y_test.values, label='Actual')
plt.plot(y_test.index, y_pred, label='Predicted')
plt.xlabel('Date')
plt.ylabel('Weekly Sales')
plt.title('Actual vs. Predicted Weekly Sales')
plt.legend()
plt.show()

```

Mean Squared Error: 303464742634.45807

